

Course: Laboratory Practice-III (Machine Learning)

Name: Anuj Mahendra Mutha

Class: BE-4

Batch : R4

Roll No. : 41443

Assignment Number : Group B - 03

Title: Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc.

Link to the Kaggle project:

<https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling>

Perform following steps:

- * Read the dataset.
- * Distinguish the feature and target set and divide the data set into training and test sets.
- * Normalize the train and test data.
- * Initialize and build the model. Identify the points of improvement and implement the same.
- * Print the accuracy score and confusion matrix (5 points).

```
In [1]: import numpy as np
        np.set_printoptions(threshold=np.inf)
        import matplotlib.pyplot as plt
        import pandas as pd
```

```
In [2]: # Importing the dataset
        dataset = pd.read_csv('/content/Churn_Modelling.csv')
```

```
In [3]: dataset.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.00
1	2	15647311	Hill	608	Spain	Female	41	1	83807.8
2	3	15619304	Onio	502	France	Female	42	8	159660.
3	4	15701354	Boni	699	France	Female	39	1	0.00
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.

```
In [4]: dataset.tail()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00
9996	9997	15569892	Johnstone	516	France	Male	35	10	5730
9997	9998	15584532	Liu	709	France	Female	36	7	0.00
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	750
9999	10000	15628319	Walker	792	France	Female	28	4	130

```
In [5]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   RowNumber           10000 non-null  int64  
 1   CustomerId          10000 non-null  int64  
 2   Surname             10000 non-null  object  
 3   CreditScore         10000 non-null  int64  
 4   Geography           10000 non-null  object  
 5   Gender              10000 non-null  object  
 6   Age                 10000 non-null  int64  
 7   Tenure              10000 non-null  int64  
 8   Balance             10000 non-null  float64 
 9   NumOfProducts       10000 non-null  int64  
10   HasCrCard           10000 non-null  int64  
11   IsActiveMember      10000 non-null  int64  
12   EstimatedSalary     10000 non-null  float64 
13   Exited              10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [6]: dataset.shape
```

```
(10000, 14)
```

```
In [7]: dataset.isnull().sum()
```

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore    0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited         0
dtype: int64
```

In [8]: `dataset.describe()`

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	Nu
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.5
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.5
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.0
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.0
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.0
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.0
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.0

In [9]: `X = dataset.iloc[:, 3:13].values`
`y = dataset.iloc[:, 13].values`

In [10]: `print(X[:10,:], '\n')`

`print(y[:10])`

```
[[619 'France' 'Female' 42 2 0.0 1 1 1 101348.88]
 [608 'Spain' 'Female' 41 1 83807.86 1 0 1 112542.58]
 [502 'France' 'Female' 42 8 159660.8 3 1 0 113931.57]
 [699 'France' 'Female' 39 1 0.0 2 0 0 93826.63]
 [850 'Spain' 'Female' 43 2 125510.82 1 1 1 79084.1]
 [645 'Spain' 'Male' 44 8 113755.78 2 1 0 149756.71]
 [822 'France' 'Male' 50 7 0.0 2 1 1 10062.8]
 [376 'Germany' 'Female' 29 4 115046.74 4 1 0 119346.88]
 [501 'France' 'Male' 44 4 142051.07 2 0 1 74940.5]
 [684 'France' 'Male' 27 2 134603.88 1 1 1 71725.73]]
```

```
[1 0 1 0 0 1 0 1 0 0]
```

```
In [13]: from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.compose import ColumnTransformer
# Country
labelencoder_X = LabelEncoder()
X[:, 1] = labelencoder_X.fit_transform(X[:, 1])

# Gender
labelencoder_X_2 = LabelEncoder()
X[:, 2] = labelencoder_X_2.fit_transform(X[:, 2])

# Giving ordinal feature to our variables
onehotencoder = ColumnTransformer([("Country", OneHotEncoder(), [1])], remainder='passthrough')
X = onehotencoder.fit_transform(X)
X = X[:, 1:]
```

```
In [14]: print(X[:10,:], '\n')
print(y[:10])

[[0.0 0.0 619 0 42 2 0.0 1 1 1 101348.88]
 [0.0 1.0 608 0 41 1 83807.86 1 0 1 112542.58]
 [0.0 0.0 502 0 42 8 159660.8 3 1 0 113931.57]
 [0.0 0.0 699 0 39 1 0.0 2 0 0 93826.63]
 [0.0 1.0 850 0 43 2 125510.82 1 1 1 79084.1]
 [0.0 1.0 645 1 44 8 113755.78 2 1 0 149756.71]
 [0.0 0.0 822 1 50 7 0.0 2 1 1 10062.8]
 [1.0 0.0 376 0 29 4 115046.74 4 1 0 119346.88]
 [0.0 0.0 501 1 44 4 142051.07 2 0 1 74940.5]
 [0.0 0.0 684 1 27 2 134603.88 1 1 1 71725.73]]

[1 0 1 0 0 1 0 1 0 0]
```

```
In [15]: X.shape

(10000, 11)
```

```
In [16]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
```

```
In [17]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [18]: import keras
import sys
from keras.models import Sequential #to initialize NN
from keras.layers import Dense #used to create layers in NN
from keras.layers import Dropout
```

```
In [20]: classifier = Sequential()
```

```
In [23]: classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation='relu'))
classifier.add(Dropout(rate=0.1))
```

```
In [24]: classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation='relu'))
classifier.add(Dropout(rate=0.1))
```

```
In [25]: classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation='sigmoid'))
```

```
In [26]: classifier.compile(optimizer = 'adam', loss= "binary_crossentropy", metrics=
```

```
In [27]: classifier.fit(X_train, y_train, batch_size = 10, epochs = 100 )
```

```
Epoch 1/100
800/800 [=====] - 3s 2ms/step - loss: 0.4854 - accuracy: 0.7960
Epoch 2/100
800/800 [=====] - 2s 2ms/step - loss: 0.4350 - accuracy: 0.7960
Epoch 3/100
800/800 [=====] - 2s 2ms/step - loss: 0.4279 - accuracy: 0.7960
Epoch 4/100
800/800 [=====] - 2s 2ms/step - loss: 0.4274 - accuracy: 0.8126
Epoch 5/100
800/800 [=====] - 2s 2ms/step - loss: 0.4272 - accuracy: 0.8260
Epoch 6/100
800/800 [=====] - 2s 2ms/step - loss: 0.4258 - accuracy: 0.8294
Epoch 7/100
800/800 [=====] - 2s 2ms/step - loss: 0.4244 - accuracy: 0.8328
```

```
In [28]: y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)

y_pred

63/63 [=====] - 0s 1ms/step
```

```
In [29]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

cm

array([[1546,  49],
       [ 274, 131]])
```

```
In [39]: from sklearn.metrics import confusion_matrix, accuracy_score, precision_score
```

```
In [40]: print("F1 Score: ", f1_score(y_test, y_pred))
```

F1 Score: 0.4478632478632479

```
In [41]: print("Accuracy: ", accuracy_score(y_test, y_pred))
```

Accuracy: 0.8385

```
In [42]: print("Precision: ",precision_score(y_test,y_pred))
        print("Recall: ",recall_score(y_test,y_pred))
```

```
Precision:  0.7277777777777777
Recall:    0.3234567901234568
```

```
In [44]: print("Error Rate: ",1-accuracy_score(y_test,y_pred))
```

```
Error Rate:  0.16149999999999998
```

```
In [32]: """Predict if the customer with the following informations will leave the
        Geography: France
        Credit Score: 600
        Gender: Male
        Age: 40
        Tenure: 3
        Balance: 60000
        Number of Products: 2
        Has Credit Card: Yes
        Is Active Member: Yes
        Estimated Salary: 50000"""
```

```
new_prediction = classifier.predict(sc.transform(np.array([[0.0, 0, 600, 1
new_prediction = (new_prediction > 0.5)
```

```
new_prediction
```

```
1/1 [=====] - 0s 19ms/step
```

```
array([[False]])
```



```
In [33]: import keras
import sys
from keras.models import Sequential #to initialize NN
from keras.layers import Dense #used to create layers in NN

#Evaluating the ANN
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import cross_val_score

def build_classifier():
    classifier = Sequential()
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu'))
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu'))
    classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
    classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])

#fit our model to the training data using KerasClassifier
classifier = KerasClassifier(build_fn = build_classifier, batch_size = 10)

#estimator - object to fit the data
#X - data to fit
#y - Target variable to try to predict
#cv - number of train test folds
#n_jobs - number of CPUs to use to do the computation. -1 means 'all CPUs'
accuracies = cross_val_score(estimator = classifier, X = X_train, y = y_train, cv = 5, n_jobs = -1)
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:18: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras (<https://github.com/adriangb/scikeras>) instead. See <https://www.adriangb.com/scikeras/stable/migration.html> (<https://www.adriangb.com/scikeras/stable/migration.html>) for help migrating.

/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:

10 fits failed out of a total of 10.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

10 fits failed with the following error:

Traceback (most recent call last):

File "/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 680, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File "/usr/local/lib/python3.7/dist-packages/keras/wrappers/scikit_learn.py", line

```
236, in fit
    return super(KerasClassifier, self).fit(x, y, **kwargs)
File "/usr/local/lib/python3.7/dist-packages/keras/wrappers/scikit_learn.py", line
157, in fit
    if (losses.is_categorical_crossentropy(self.model.loss) and
AttributeError: 'NoneType' object has no attribute 'loss'

warnings.warn(some_fits_failed_message, FitFailedWarning)
```